

Name: Gore Aniket Machhindra

Email: [aniket.m.gore.1901@gmail.com](mailto:aniket.m.gore.1901@gmail.com)

Role: Machine Learning Intern

## Task 4: Housing Prices Prediction Project

### Content -

- Importing Libraries
- Importing the Dataset
- Initializing the DataFrame
- Adding Feature Names to the dataframe
- Adding target variable to dataframe
- Checking Missing Values
- Describe the Data
- Correlation
- Heatmap
- Splitting the Data
- Importing random Forest regressor
- Model Prediction on Training Data
- Model Evaluation
- Visualizing the differences between actual prices and predicted values
- Conclusion

### ▼ Importing Libraries

```
1 import pandas as pd
2 import numpy as np
3 from sklearn import metrics
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 %matplotlib inline
```

## ▾ Importing the Dataset

```
1 from sklearn.datasets import load_boston
2 boston = load_boston()
```

## ▾ Initializing the Dataframe

```
1 data = pd.DataFrame(boston.data)
```

```
1 data.head()
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	

## ▾ Adding Feature Names to the dataframe

```
1 data.columns = boston.feature_names
2 data.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	

## ▾ Adding target variable to dataframe

```
1 data['PRICE'] = boston.target
```

```
1 data.shape
```

```
(506, 14)
```

```
1 data.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
      'PTRATIO', 'B', 'LSTAT', 'PRICE'],  
      dtype='object')
```

```
1 data.dtypes
```

```
CRIM      float64  
ZN        float64  
INDUS     float64  
CHAS      float64  
NOX       float64  
RM        float64  
AGE       float64  
DIS       float64  
RAD       float64  
TAX       float64  
PTRATIO   float64  
B         float64  
LSTAT     float64  
PRICE     float64  
dtype: object
```

## ▼ Checking Missing Values

```
1 data.isnull().sum()
```

```
CRIM      0  
ZN        0  
INDUS     0  
CHAS      0  
NOX       0  
RM        0  
AGE       0  
DIS       0  
RAD       0  
TAX       0  
PTRATIO   0  
B         0  
LSTAT     0
```

```
PRICE      0
dtype: int64
```

## Describe the Data

```
1 data.describe()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
<b>mean</b>	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
<b>std</b>	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
<b>min</b>	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
<b>25%</b>	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
<b>50%</b>	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
<b>75%</b>	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
<b>max</b>	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000



## Correlation


```
1 corr = data.corr()
2 corr.shape
```

```
(14, 14)
```

## Heatmap

```
1 plt.figure(figsize=(20,20))
2 sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':15}, c
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2fb7570b90>



## ▼ Splitting the Data

```
1 X = data.drop(['PRICE'], axis = 1)
2 y = data['PRICE']
```

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3, random_state = 4
```

## ▼ Importing random Forest regressor

```
1 from sklearn.ensemble import RandomForestRegressor
2 reg = RandomForestRegressor()
3 reg.fit(X_train, y_train)
```

```
RandomForestRegressor()
```

## ▼ Model Prediction on Training Data

```
1 y_pred = reg.predict(X_train)
```

## ▼ Model Evaluation

```
1 print('R^2:',metrics.r2_score(y_train, y_pred))
```

```
R^2: 0.9786402541948575
```

```
1 print('Adjusted R^2:',1 - (1-metrics.r2_score(y_train, y_pred))*(len(y_train)-1)/(len(y_train)-2))
```

```
Adjusted R^2: 0.9778235580317196
```

```
1 print('MAE:',metrics.mean_absolute_error(y_train, y_pred))
2 print('MSE:',metrics.mean_squared_error(y_train, y_pred))
3 print('RMSE:',np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```


MAE: 0.8485790960451979  
MSE: 1.6077657881355942  
RMSE: 1.2679770455870225

## Visualizing the differences between actual prices and predicted values

```
1 plt.scatter(y_train, y_pred)
2 plt.xlabel("Prices")
3 plt.ylabel("Predicted prices")
4 plt.title("Prices vs Predicted prices")
5 plt.show()
```



```
1 plt.scatter(y_pred, y_train-y_pred)
2 plt.title("Predicted vs residuals")
3 plt.xlabel("Predicted")
4 plt.ylabel("Residuals")
5 plt.show()
6
```



## ▼ Conclusion

- Random Forest regressor is works best for this dataset
- R2 Score is 97% Accurate on this Dataset
- Adjusted R2 Score is 97% Accurate on this Dataset